

Sales Pulse: Sales and Demand Forecasting System

Y. Nimisha

*Department of Computer Science
and Engineering
Koneru Lakshmaiah Education
Foundation
Hyderabad, Telangana, India
nimishayarlagadda@gmail.com*

K. Sanuhya

*Department of Computer
Science and Engineering
Koneru Lakshmaiah
Education Foundation
Hyderabad, Telangana,
India
sanuhya5@gmail.com*

R. Sarvajith Varma

*Department of Computer
Science and Engineering
Koneru Lakshmaiah Education
Foundation
Hyderabad, Telangana, India
sarvajithvarma2007@gmail.com*

Dr. N. Sunanda

*Department of Computer
Science and Engineering
Koneru Lakshmaiah
Education Foundation
Hyderabad, Telangana,
India*

Abstract—Businesses generate large amounts of historical sales data through their daily transactions. Analyzing this data is important for understanding sales trends, demand variations, and seasonal patterns. However, traditional sales management and reporting systems primarily focus on recording transactions and presenting past sales information, with limited ability to predict future sales and demand. This can make it difficult for businesses to identify high- and low-demand periods and plan inventory effectively. This paper presents Sales Pulse, a sales and demand forecasting system designed to transform historical sales data into future sales predictions and useful business insights. The proposed system involves collecting and preprocessing historical sales data, performing exploratory data analysis to identify trends and seasonal variations, and applying time-series forecasting techniques such as ARIMA or Prophet. Forecasting performance can be evaluated using suitable accuracy metrics. The system is also intended to provide a user-friendly dashboard for displaying historical sales, forecasted values, and important insights. Sales Pulse aims to support proactive decision-making, improve inventory planning, and provide businesses with a better understanding of their sales patterns.

Keywords—Sales forecasting, demand forecasting, time-series forecasting, sales analysis, ARIMA, Prophet, inventory planning, data science.

I. INTRODUCTION

Sales data plays an important role in understanding business performance and customer demand. Businesses can use historical sales records to identify changes in demand, recognize seasonal patterns, and understand periods of high and low sales. Effective analysis of such information can support better planning and decision-making.

Traditional sales management and reporting systems generally focus on recording transactions and displaying past sales information. Although these systems are useful for monitoring historical performance, they provide limited support for predicting future sales and demand. As a result, businesses may face difficulties in anticipating future requirements and planning inventory efficiently.

Time-series forecasting provides a data-driven approach to addressing this limitation. By analyzing historical observations over time, forecasting techniques can be used to estimate future values and identify patterns in demand. In this context, Sales Pulse is proposed as a sales and demand forecasting system that combines historical data analysis with time-series forecasting.

The main purpose of Sales Pulse is to transform historical sales data into useful future sales predictions and business insights. The system focuses on data collection,

preprocessing, exploratory analysis, forecasting, performance evaluation, and visualization. By presenting

historical and forecasted sales information through a user-friendly dashboard, the system aims to support informed and proactive business decisions.

II. PROBLEM STATEMENT

Businesses need to analyze historical sales data to understand sales trends and demand variations. Traditional sales management and reporting systems mainly record transactions and display past sales information. However, these systems have limited ability to predict future sales and demand.

This limitation can make it difficult for businesses to identify high-demand and low-demand periods and plan inventory effectively. Without appropriate forecasting, businesses may find it challenging to anticipate changes in demand and make proactive decisions.

Sales Pulse aims to address this problem by using data science and time-series forecasting to predict future sales and provide useful business insights.

III. OBJECTIVES

The major objectives of the Sales Pulse system are to collect and analyze historical sales data; preprocess data by handling missing, inconsistent, and incorrect records; identify sales trends, seasonal patterns, and demand variations through exploratory data analysis; apply time-series forecasting models such as ARIMA or Prophet to predict future sales and demand; evaluate forecasting performance using suitable accuracy metrics; develop a user-friendly dashboard to display historical sales, forecasted values, and important insights; and support better inventory planning and informed business decision-making.

IV. LITERATURE SURVEY

Existing research demonstrates the application of different machine learning and time-series forecasting techniques for sales and demand prediction.

A comparative study on time-series forecasting of seasonal item sales investigated forecasting approaches using furniture sales data. While the study provides useful insights into forecasting seasonal products, its focus on a particular product category may limit the generalization of its results to other businesses and product categories [1].

Hasan, Kabir, Shuvro, and Das presented a comparative study of retail sales forecasting using Walmart historical sales data. The study provides a basis for comparing forecasting approaches, although model performance may vary when applied to different businesses and datasets [2].

Another study investigated predictive analytics for demand forecasting by comparing SARIMA and LSTM

approaches in retail supply chain management. The study used sales data from a specific Austrian retailer, indicating that forecasting performance can depend on product demand patterns and promotional factors [3].

Brykin compared ARIMA, LSTM, and Prophet models for sales forecasting. The study demonstrates the use of different forecasting approaches, but the best-performing model may vary depending on the sales category, dataset, and forecasting period [4].

Song, Du, and Li investigated precise demand forecasting using a joint LSTM-ARIMA model. Although combining forecasting approaches can improve modeling capability, joint models may also increase computational complexity and implementation requirements [5].

Ecevit, Öztürk, Dağ, and Özcan investigated short-term sales forecasting using LSTM and Prophet-based approaches in e-commerce. The study focuses primarily on short-term e-commerce forecasting, and therefore the results may not necessarily generalize to long-term forecasting or other industries [6].

These studies indicate that forecasting performance is influenced by the characteristics of the dataset, product demand patterns, forecasting period, and selected model. Based on these observations, Sales Pulse focuses on analyzing historical sales data and applying time-series forecasting techniques such as ARIMA or Prophet to generate future sales predictions.

V. PROPOSED SYSTEM

Sales Pulse is proposed as a data science-based sales and demand forecasting system. The system transforms historical sales records into useful predictions and business insights.

VI. SYSTEM METHODOLOGY

The Sales Pulse methodology is organized as a sequential data-to-insight pipeline. Each stage prepares the output required by the next stage, while the evaluation stage provides a mechanism for assessing whether the generated forecasts are sufficiently useful for decision support. The methodology is deliberately modular so that preprocessing, exploratory analysis, forecasting, evaluation, and visualization can be improved independently without changing the complete workflow.

A. Historical Sales Data Collection

The first stage involves collecting historical sales records that describe how products or categories have performed over time. The historical dataset forms the foundation of the forecasting process because the forecasting models learn temporal behavior from previously observed values. Depending on the business context, the records may contain dates, sales quantities, revenue-related measures, product identifiers, categories, or other fields that help describe demand. The proposed system treats time as an essential dimension and therefore organizes the usable observations into an ordered sequence before forecasting.

Data collection also establishes the observation period and the forecasting horizon. A consistent time index is important because irregular or incorrectly ordered observations can affect the interpretation of trends and seasonal behavior. At this stage, the system therefore focuses on obtaining a

reliable historical series rather than immediately applying a prediction model.

B. Data Preprocessing

Raw transaction data may contain missing values, inconsistent formats, duplicate records, incorrect entries, or other quality problems. Preprocessing is therefore performed before exploratory analysis and forecasting. The objective is not simply to remove records, but to create a consistent analytical dataset in which dates and sales measures can be interpreted correctly.

The preprocessing stage can include checking the date field, identifying missing observations, examining duplicate or inconsistent records, and preparing the sales measure for time-series analysis. Records that cannot be interpreted reliably are identified and handled before model training. The resulting dataset is arranged chronologically so that subsequent analysis preserves the temporal relationship between observations.

C. Exploratory Data Analysis

Exploratory data analysis (EDA) is used to understand the behavior of the prepared sales series before selecting or applying a forecasting model. The proposed analysis focuses on overall sales trends, changes in demand, high- and low-demand periods, and possible seasonal variations. Visualization is especially useful at this stage because temporal patterns that are difficult to recognize from individual records can become visible when sales are plotted over time.

Trend analysis helps determine whether sales generally increase, decrease, or remain relatively stable over the observation period. Seasonal analysis examines whether recurring patterns appear at particular periods. Demand variation analysis helps identify periods in which sales differ substantially from the surrounding observations. These observations provide contextual information for the forecasting stage and can also help explain forecast behavior to end users.

D. Time-Series Forecasting

After preprocessing and EDA, the system applies time-series forecasting techniques to estimate future sales and demand. The proposed paper identifies ARIMA and Prophet as candidate forecasting approaches. ARIMA represents a classical statistical time-series approach, while Prophet provides a model designed for time-dependent forecasting with interpretable trend and seasonal components. The final model selection is intended to depend on the characteristics of the available sales data and the forecasting task.

The forecasting stage uses historical observations as the input and produces predicted values for a defined future horizon. Forecast generation should preserve the temporal ordering of the data so that information from future observations is not unintentionally used when estimating earlier periods. This temporal separation is important for obtaining a meaningful assessment of forecasting performance.

E. Forecast Evaluation

Forecast evaluation is included to determine how closely predicted sales correspond to observed sales. Suitable accuracy metrics can be used to quantify forecasting error. The paper does not prescribe a single metric as universally

optimal; instead, the evaluation stage is intended to provide numerical evidence that can be interpreted together with the characteristics of the sales dataset.

Evaluation can be organized using historical observations for model development and a later portion of the series for testing. Comparing predictions against known observations provides an empirical basis for assessing the usefulness of a forecasting approach. The evaluation results can also support model comparison when more than one candidate model is implemented.

F. Dashboard and Visualization

The final stage presents the results through a user-friendly dashboard. The dashboard is intended to combine historical sales information, forecasted values, and important insights in a form that can be understood without directly examining the underlying analytical data. A visual comparison of historical and predicted values can help users recognize the direction of demand and identify periods requiring attention. The dashboard also acts as the communication layer of the system. Instead of exposing only model outputs, it connects forecasts with the patterns discovered during EDA and the evaluation results. This supports a workflow in which data is transformed into information and then presented as business-oriented insight. The visualization layer is therefore an important part of Sales Pulse rather than an optional reporting component.

G. System Architecture



Fig. 1. System architecture of the proposed Sales Pulse sales and demand forecasting system.

Figure 1 summarizes the logical architecture of Sales Pulse. Historical sales data enters the preprocessing layer, where data quality and temporal consistency are addressed. The prepared series is then examined through EDA to identify trends, seasonal behavior, and demand variations. The forecasting layer applies candidate time-series models such as ARIMA or Prophet. Forecast evaluation provides a quantitative assessment of prediction quality, after which the results are exposed through the dashboard. The final insight layer connects the analytical output to inventory and sales planning activities.

The architecture separates data preparation, analysis, prediction, evaluation, and presentation into distinct logical

stages. This separation supports maintainability because changes to one stage can be made without redesigning the entire workflow. It also makes the system easier to explain to users and evaluators because the path from historical records to business insights is explicit.

H. End-to-End Processing

In an end-to-end execution, the system first receives historical sales information and validates the available records. The cleaned series is then explored to establish an understanding of its temporal behavior. Based on this analysis, the forecasting stage generates future values using a selected candidate model. The predicted values are evaluated against appropriate historical observations, and the final results are prepared for visualization.

The sequence is intentionally designed so that forecasting is not treated as an isolated machine-learning operation. Data quality and exploratory understanding precede model application, while evaluation follows prediction. The dashboard then communicates both historical and future-oriented information. This structure supports the central objective of Sales Pulse: converting historical sales data into a form that can assist proactive planning.

I. Data Quality and Reliability Considerations

The reliability of a forecasting system depends strongly on the quality and representativeness of the historical data supplied to it. Missing or inconsistent records can affect trend estimation, while unusual observations may influence model behavior. For this reason, Sales Pulse includes preprocessing and EDA as explicit stages rather than treating them as minor preparation tasks.

Another consideration is that forecasting accuracy can vary with the product category, demand pattern, forecasting horizon, and dataset characteristics. The literature surveyed in this paper reports differences across forecasting approaches and application contexts. Consequently, the proposed system emphasizes evaluation rather than assuming that one forecasting model will always provide the most useful predictions.

J. Model Selection Considerations

ARIMA and Prophet are included as candidate approaches because the proposed system is centered on time-series forecasting. Their outputs can be examined using the same evaluation framework, allowing model behavior to be compared for the available sales series. If one model captures the observed temporal behavior more effectively, its forecasts can be used for the dashboard and downstream interpretation.

The proposed methodology also leaves room for future extension. Additional forecasting approaches can be incorporated without changing the overall architecture, provided that they accept the prepared time-series data and return predictions that can be evaluated and visualized. This modularity is useful when the characteristics of the business dataset change over time.

VII. SYSTEM WORKFLOW

The overall workflow of Sales Pulse can be represented as: Historical Sales Data → Data Preprocessing → Exploratory Data Analysis → Time-Series Forecasting → Forecast Evaluation → Dashboard and Insights.

The system begins with historical sales data. The data is cleaned and prepared during preprocessing. Exploratory analysis is then used to identify trends, seasonal patterns, and demand variations. Forecasting models such as ARIMA or Prophet are subsequently applied to predict future sales and demand. The forecasting results are evaluated using suitable accuracy metrics. Finally, the historical data, predictions, and important insights are presented through a dashboard.

VIII. EXPECTED OUTCOMES

The expected outcomes of Sales Pulse include accurate and useful future sales predictions; clear visualization of historical sales trends and forecasted demand; identification of high-demand and low-demand periods; better understanding of sales patterns and demand variations; improved support for inventory and sales planning; and better support for informed and proactive business decision-making.

IX. REAL-WORLD RELEVANCE

Sales forecasting can be useful for businesses that need to plan inventory and sales activities based on expected demand. By identifying seasonal patterns and demand variations, businesses can gain a better understanding of how sales change over time.

Sales Pulse is intended to support this process by converting historical sales information into future predictions and visual insights. The availability of forecasted demand can help businesses make more informed decisions regarding inventory and sales planning.

From a practical perspective, Sales Pulse is intended for settings in which historical sales records are available and future demand needs to be estimated. The system does not replace operational planning; rather, it provides an analytical view that can be considered alongside business knowledge and other operational constraints. Historical and forecasted information can help users examine whether expected demand is increasing, decreasing, or showing recurring variation.

Inventory planning is one of the principal areas in which demand forecasts can be useful. When a business has an estimate of future demand, planning activities can be organized around expected high- and low-demand periods. The usefulness of the forecast depends on data quality and model performance, so forecast values should be interpreted together with evaluation results and the underlying historical pattern.

Sales Pulse can also support communication between technical and non-technical users. Analysts can use the EDA and evaluation stages to examine model behavior, while decision-makers can use the dashboard to view summarized historical and forecast information. This separation between analytical processing and visualization makes the proposed system suitable as a practical decision-support prototype.

An additional benefit of the proposed architecture is traceability. Each forecast can be related back to the historical data, preprocessing stage, exploratory findings, forecasting model, and evaluation stage that produced it. Such traceability can help users understand that a forecast is the output of a sequence of analytical operations rather than an unexplained numerical value.

A. Operational Decision Support

Sales Pulse is designed as a decision-support system rather than as an autonomous decision-making mechanism. The forecast provides an estimate of future sales based on historical observations, while the final operational decision remains dependent on business context. Factors such as stock availability, supplier constraints, pricing changes, promotions, and unexpected events can influence actual demand and may not be represented in a historical sales series.

The dashboard therefore has an important role in presenting the forecast together with the historical context from which it was generated. A user can examine the direction of historical sales, observe forecasted values, and identify periods that may require additional planning attention. This approach is consistent with the objective of transforming raw historical information into accessible analytical evidence.

B. Inventory Planning Use Case

Inventory planning is a direct application area for demand forecasting. When historical data indicates recurring periods of increased demand, the forecast can help users anticipate the possibility of higher future requirements. Similarly, periods of lower expected demand can be identified for planning purposes. The proposed system does not prescribe a fixed inventory quantity; instead, it supplies forecast information that can be incorporated into an organization's existing planning process.

The value of this approach is particularly related to timing. Historical reporting describes what has already occurred, whereas forecasting introduces a forward-looking component. By presenting both views on the same dashboard, Sales Pulse can help users compare past behavior with expected future demand and use that comparison when organizing sales and inventory activities.

C. Sales Planning Use Case

Sales teams can also use forecast information to understand expected demand patterns. The identification of high-demand and low-demand periods can assist with planning activities that are sensitive to timing. The system can make these patterns easier to communicate because the dashboard provides a visual representation instead of requiring users to inspect individual historical records.

Because the paper proposes a general forecasting workflow rather than a single industry-specific deployment, the system can be interpreted as a reusable analytical framework. Different businesses may supply different sales datasets, but the major processing stages remain applicable: collection, cleaning, exploration, forecasting, evaluation, and visualization.

D. Interpretability and User Interaction

Interpretability is important when forecast outputs are used by people who are not directly involved in model development. Sales Pulse addresses this requirement by placing forecasting inside a broader analytical workflow. The dashboard is intended to expose historical trends and forecasted values in an understandable form, while the EDA stage provides additional context for observed patterns.

The proposed architecture also separates model computation from result presentation. This separation means that users

can interact with the final analytical output without needing to understand the internal details of ARIMA or Prophet. At the same time, technical users can inspect the preprocessing and evaluation stages when deeper analysis is required.

E. Scalability of the Proposed Workflow

The logical architecture is designed so that each processing stage can be extended as the dataset or business requirement grows. For example, additional preprocessing rules can be introduced when new data-quality issues are observed, while additional forecasting models can be evaluated without changing the purpose of the dashboard. This modular organization is useful for a system that may evolve from a prototype into a broader analytical application.

Scalability in this context should be understood at both the data and functional levels. At the data level, the pipeline can process a larger historical record if the implementation environment supports it. At the functional level, new analytical views, forecast horizons, or evaluation procedures can be incorporated while retaining the central data-to-insight workflow.

F. Security, Reliability, and Data Governance Considerations

Although the proposed paper primarily focuses on forecasting, a practical deployment would also require attention to data governance. Historical sales information can represent commercially important business records, so access to source data and dashboard outputs should be controlled according to organizational requirements. The forecasting pipeline should also preserve the distinction between raw records and processed analytical data so that changes can be traced.

Reliability is similarly connected to repeatability. A consistent preprocessing procedure and a clearly defined forecasting horizon help ensure that results can be regenerated from the same analytical inputs. Recording the model used, the evaluation approach, and the period represented by the forecast can further improve the transparency of the system.

G. Evaluation Strategy

The proposed evaluation strategy should consider both numerical forecasting accuracy and practical usefulness. Numerical evaluation provides a measurable indication of prediction error, while visual inspection can reveal whether the model captures broad trends or recurring variations. Using both forms of assessment is appropriate because a single numerical value may not fully describe the behavior of a time-series forecast.

A suitable evaluation process can reserve later historical observations as a test period and compare the model's predictions with the known values. Candidate models such as ARIMA and Prophet can then be examined using the same evaluation procedure. This provides a consistent basis for understanding how different approaches behave on the available dataset without assuming in advance that one approach will always perform better.

Evaluation should also be interpreted in relation to the forecast horizon. A model that is useful for a short horizon may behave differently when predictions are extended further into the future. Therefore, the forecasting period

should be defined clearly when reporting system performance.

H. Limitations

The current proposal has several limitations. First, the paper does not present a completed empirical comparison with measured forecast accuracy; ARIMA and Prophet are described as candidate models. Second, the available discussion is centered on historical sales variables and does not explicitly incorporate external drivers such as promotions, economic conditions, competitor activity, or weather. Such variables can affect real-world demand and may influence forecasting performance.

Third, the usefulness of the proposed system depends on the quality and length of the historical dataset. Sparse records, missing observations, abrupt structural changes, or highly irregular demand may make forecasting more difficult. These limitations do not invalidate the proposed architecture, but they indicate that model results should be interpreted in the context of the data used.

I. Future Scope

Future development can extend Sales Pulse in several directions. Additional forecasting methods can be incorporated and compared within the same evaluation framework. More detailed product-level or category-level forecasting can also be explored when the available data supports such analysis. The dashboard can be extended with interactive filtering, forecast-horizon selection, and richer comparisons between historical and predicted values.

Another possible extension is the inclusion of external variables and more advanced hybrid approaches. The literature surveyed in this paper includes work involving LSTM, SARIMA, and joint LSTM-ARIMA approaches, indicating that forecasting systems can be expanded beyond the initial candidate models. Such extensions should be evaluated empirically on the target dataset rather than assumed to improve performance automatically.

Future work can also focus on automated model selection and periodic retraining when new sales observations become available. A continuously updated pipeline would allow the system to incorporate recent demand behavior while retaining the same high-level architecture. Such an extension would move Sales Pulse closer to a production-oriented forecasting platform.

VII. SYSTEM WORKFLOW

A. Workflow Interpretation

The workflow of Sales Pulse can be interpreted as a sequence of transformations in which each stage has a defined analytical purpose. Historical sales records are first converted into a consistent dataset. The cleaned dataset is then transformed into descriptive information through EDA, after which the forecasting stage transforms historical observations into estimates of future demand. Evaluation transforms forecast results into measurable evidence about model performance, and the dashboard transforms that evidence into information that can be consumed by users.

The sequential arrangement reduces the risk of treating forecasting as an isolated modeling task. A forecast is meaningful only when its input data is reliable and when its

performance can be assessed. For this reason, the workflow places preprocessing and evaluation around the forecasting component. The dashboard is positioned after these analytical stages so that users receive results that have passed through the complete proposed pipeline.

B. Input and Output of Each Stage

The input to the collection stage is historical sales information, while its output is an ordered dataset suitable for analysis. The preprocessing stage receives this dataset and produces a cleaned analytical series. EDA receives the prepared series and produces observations about trends, seasonal patterns, and demand variation. The forecasting stage uses the prepared time series to produce future estimates. Evaluation receives the predicted values and corresponding observed values used for testing and produces accuracy information. Finally, the dashboard receives the analytical results and presents them visually.

Defining these inputs and outputs makes the system easier to implement and test. If an error occurs in the final dashboard, the logical architecture allows the implementation to be traced backward through the forecasting, evaluation, EDA, and preprocessing stages. This staged structure can therefore support both development and maintenance of the proposed system.

C. Information Flow and Feedback

Although the main workflow is sequential, the evaluation results can provide feedback for later model selection and refinement. If a candidate forecasting approach does not represent the observed series adequately, the model can be reconsidered and the forecasting stage repeated. Similarly, EDA can reveal data characteristics that require an adjustment in preprocessing before forecasting is performed again.

This feedback relationship is important because forecasting is an iterative analytical process. The architecture therefore represents a logical pipeline rather than a single irreversible execution. Historical data remains the foundation, while evaluation and exploratory findings can guide subsequent iterations of the analytical workflow.

VIII. EXPECTED OUTCOMES

A. Analytical Outcomes

The analytical outcome of Sales Pulse is a structured representation of historical and future-oriented sales behavior. The system is expected to identify meaningful trends, seasonal variations, and demand differences and to convert these observations into forecasted values. The expected result is not only a numerical prediction but also a clearer understanding of how sales change over time. The forecasting output can be represented alongside historical observations so that the difference between observed and expected future behavior is visible. Such visualization can make the output easier to interpret and can support the identification of periods that may require additional planning attention.

B. Business-Oriented Outcomes

At the business level, the proposed system is expected to support inventory planning, sales planning, and proactive

decision-making. These outcomes follow directly from the paper's central objective of transforming historical records into future sales predictions and useful insights. The system does not claim that forecasting alone guarantees improved business performance; rather, it provides information that can be incorporated into planning processes.

The dashboard is expected to make these outputs accessible to users by presenting historical sales, forecasted values, and important insights in a single analytical view. A unified presentation can reduce the need to interpret multiple disconnected reports and can help users compare historical behavior with expected future demand.

C. Evaluation Outcomes

A further expected outcome is a measurable understanding of forecasting performance. By applying suitable accuracy metrics, the system can indicate the degree to which predictions correspond to observed values during evaluation. This provides a quantitative basis for assessing the usefulness of the forecasting stage and for comparing candidate approaches when multiple models are implemented.

The evaluation results should be interpreted with the characteristics of the dataset and forecasting horizon in mind. A numerical metric is useful for comparison, but it does not remove the need for data-quality checks and visual analysis. Combining these perspectives is consistent with the overall design of Sales Pulse as an integrated analytical system.

X. CONCLUSION

Sales Pulse presents a data science-based approach to sales and demand forecasting. Traditional sales reporting systems mainly provide historical information, which may not be sufficient for businesses that need to anticipate future demand. The proposed system addresses this limitation by combining historical sales analysis, data preprocessing, exploratory data analysis, and time-series forecasting.

Models such as ARIMA and Prophet can be applied to generate future sales and demand predictions, while suitable accuracy metrics can be used to evaluate their performance. A user-friendly dashboard can further improve the accessibility of historical and forecasted information.

The proposed system is expected to provide useful sales predictions, identify high- and low-demand periods, improve understanding of sales patterns, and support inventory and sales planning. Overall, Sales Pulse aims to transform historical sales data into actionable information that can support informed and proactive business decisions.

REFERENCES

- [1] "Time-series forecasting of seasonal items sales using machine learning – A comparative analysis," *Journal of Innovation & Knowledge*, 2022, doi: 10.1016/j.jjime.2022.100058.
- [2] M. R. Hasan, M. A. Kabir, R. A. Shuvro, and P. Das, "A comparative study on forecasting of retail sales," *arXiv preprint arXiv:2203.06848*, 2022.
- [3] "Predictive analytics for demand forecasting – A comparison of SARIMA and LSTM in retail SCM," *Procedia Computer Science*, 2022, doi: 10.1016/j.procs.2022.01.298.
- [4] D. Brykin, "Sales forecasting models: Comparison between ARIMA, LSTM and Prophet," *Journal of*

Computer Science, 2024, doi:
10.3844/jcssp.2024.1222.1230.

- [5] H. Song, D. Du, and Y. Li, "Research on precise demand forecast of retail commodities based on LSTM-ARIMA joint model," *Journal of Physics: Conference Series*, vol. 1865, no. 4, 2021, doi: 10.1088/1742-6596/1865/4/042049.
- [6] A. Ecevit, İ. Öztürk, M. Dağ, and T. Özcan, "Short-term sales forecasting using LSTM and Prophet based models in e-commerce," 2023.